

Computer Vision I (CSE 40535/60535)

Practical 4: CNN-based object detection

Friday, November 15, 2019

Tasks to do at home **before class**

To work on practical 4, we are going to need Python3 (for instance 3.6.x) and a few libraries installed, in addition to the OpenCV library. **Please do not use Python 3.7.x for this practical.**

1. **If you installed OpenCV and Python3 (<3.7) to complete task 1 (color-based object detection), you can skip this step.** If you don't have Python3 and OpenCV, you can use your preferred way to install them, for instance install Anaconda Python (<https://www.anaconda.com/download/>), and then use its package manager to install OpenCV:

```
$ conda install -c conda-forge opencv
```

2. Now, all you need to do is to install scikit-learn, tensorflow and keras. This can be done with these commands:

If you are using pip:

```
$ pip install numpy scipy scikit-learn  
$ pip install h5py  
$ pip install tensorflow  
$ pip install keras
```

If you are using anaconda:

```
$ conda install -c anaconda numpy  
$ conda install -c anaconda scipy  
$ conda install -c anaconda scikit-learn  
$ conda install -c anaconda h5py  
$ conda install -c conda-forge tensorflow  
$ conda install -c conda-forge keras
```

If you encountered any problems during these steps, these sites can help:

<http://scikit-learn.org/stable/install.html>

<https://www.tensorflow.org/install/>

<https://keras.io/#installation>

3. Before using keras, we have to make sure it is set up to use tensorflow as the backend. Follow these instructions to make sure it is correctly configured: <https://keras.io/backend/>
4. Test the keras installation by calling Python and importing the keras module:

```
$ python  
>>> import keras
```

You should see a message saying that keras is using tensorflow backend.

5. Run the example code:

```
$ python main.py --classifier resnet
```

The first time the program runs, it will automatically download the trained weights for ResNet Convolutional Neural Network [<https://arxiv.org/abs/1512.03385>]. This program will use a stream of video from your webcam as the input and use ResNet to classify objects in the image.



You should be able to see the top 3 classes and the confidence associated with them in upper left corner. For this example, they are:

Computer keyboard: 0.966

Spade bar: 0.007

Laptop: 0.007

Here you can check the names and images of 1000 classes which the ResNet and VGG were trained for:

<http://image-net.org/challenges/LSVRC/2015/browse-synsets>

Present to the camera objects that belong to the classes used in training. For instance “American alligator” ☺

6. If you run the program using `--classifier vgg`, the VGG Convolutional Neural Network [<https://arxiv.org/abs/1409.1556>] will be downloaded and used:



You should be able to see the top-3 classes and the probabilities for them on upper left corner. For this example, they are:

Computer keyboard: 0.130

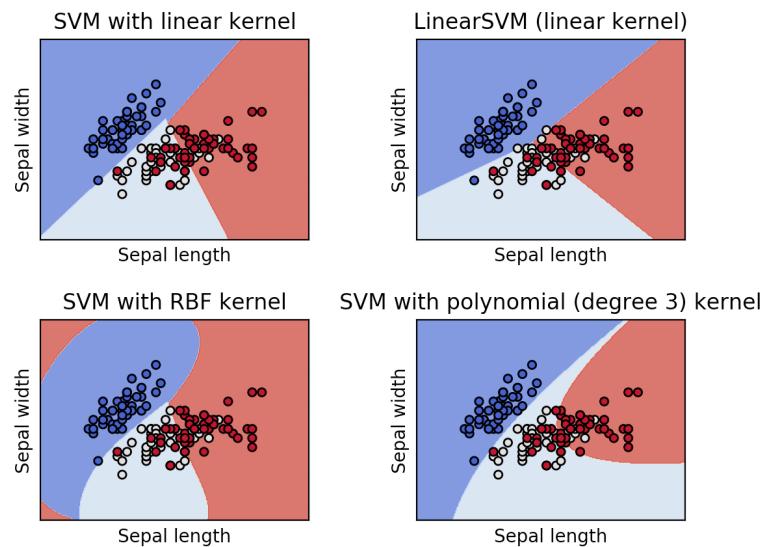
Notebook: 0.128

Mouse: 0.073

7. Alternatively, you can run the program without any parameters, and press ‘r’ or ‘v’ to capture a frame and classify it using respectively ResNet or VGG.
8. Just to double check that the `scikit-learn` and `matplotlib` are up and running, run this example code that illustrates the SVM-based multiclass classification on a very simple dataset of flower features:

```
$ python svm_plot_iris.py
```

If everything works well, you should see the following picture:



Trouble shooting:

If you are using MacOS and encounter problems with matplotlib backend when running this example, you can find the code at this URL: https://scikit-learn.org/0.18/auto_examples/svm/plot_iris.html

You may copy and paste the code to a separate python script, and add following lines at the beginning of the script to set the backend:

```
from sys import platform as sys_pf
if sys_pf == 'darwin':
    import matplotlib
    matplotlib.use("TkAgg")
import matplotlib.pyplot as plt
plt.plot()
```